

ML algorithms

November 1, 2020

```
[ ]: #Linear Regression
```

```
[2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import pylab as pl
%matplotlib inline
```

```
[3]: df = pd.read_csv('FuelConsumptionCo2.csv')
```

```
[4]: df.head()
```

```
[4]:
```

	MODELYEAR	MAKE	MODEL	VEHICLECLASS	ENGINE SIZE	CYLINDERS	\
0	2014	ACURA	ILX	COMPACT	2.0	4	
1	2014	ACURA	ILX	COMPACT	2.4	4	
2	2014	ACURA	ILX HYBRID	COMPACT	1.5	4	
3	2014	ACURA	MDX 4WD	SUV - SMALL	3.5	6	
4	2014	ACURA	RDX AWD	SUV - SMALL	3.5	6	

	TRANSMISSION	FUELTYPE	FUELCONSUMPTION_CITY	FUELCONSUMPTION_HWY	\
0	AS5	Z	9.9	6.7	
1	M6	Z	11.2	7.7	
2	AV7	Z	6.0	5.8	
3	AS6	Z	12.7	9.1	
4	AS6	Z	12.1	8.7	

	FUELCONSUMPTION_COMB	FUELCONSUMPTION_COMB_MPG	CO2EMISSIONS
0	8.5	33	196
1	9.6	29	221
2	5.9	48	136
3	11.1	25	255
4	10.6	27	244

```
[5]: df.describe()
```

```
[5]:
```

	MODELYEAR	ENGINE SIZE	CYLINDERS	FUELCONSUMPTION_CITY	\
count	1067.0	1067.000000	1067.000000	1067.000000	
mean	2014.0	3.346298	5.794752	13.296532	

std	0.0	1.415895	1.797447	4.101253
min	2014.0	1.000000	3.000000	4.600000
25%	2014.0	2.000000	4.000000	10.250000
50%	2014.0	3.400000	6.000000	12.600000
75%	2014.0	4.300000	8.000000	15.550000
max	2014.0	8.400000	12.000000	30.200000


	FUELCONSUMPTION_HWY	FUELCONSUMPTION_COMB	FUELCONSUMPTION_COMB_MPG	\
count	1067.000000	1067.000000	1067.000000	
mean	9.474602	11.580881	26.441425	
std	2.794510	3.485595	7.468702	
min	4.900000	4.700000	11.000000	
25%	7.500000	9.000000	21.000000	
50%	8.800000	10.900000	26.000000	
75%	10.850000	13.350000	31.000000	
max	20.500000	25.800000	60.000000	

	CO2EMISSIONS
count	1067.000000
mean	256.228679
std	63.372304
min	108.000000
25%	207.000000
50%	251.000000
75%	294.000000
max	488.000000


```
[6]: #simple linear regression
```

```
[7]: cdf = df[['ENGINE_SIZE', 'CYLINDERS', 'FUELCONSUMPTION_COMB', 'CO2EMISSIONS']]
```


```
[8]: viz = cdf
viz.hist()
plt.show()
```




```
[9]: plt.scatter(cdf.FUELCONSUMPTION_COMB, cdf.CO2EMISSIONS, color='blue')
plt.xlabel('FUELCONSUMPTION')
plt.ylabel('emission')
plt.show()
```



```
[10]: plt.scatter(cdf.ENGINESIZE, cdf.CO2EMISSIONS)
plt.xlabel('Engine size')
plt.ylabel('emission')
plt.show()
```




```
[11]: plt.scatter(cdf.CYLINDERS, cdf.CO2EMISSIONS)
plt.xlabel('no. of cylinders')
plt.ylabel('emission')
plt.show()
```



```
[12]: msk = np.random.rand(len(df)) < 0.8  
train = cdf[msk]  
test = cdf[~msk]
```

```
[13]: plt.scatter(train.ENGINESIZE, train.CO2EMISSIONS)  
plt.xlabel('engine')  
plt.ylabel('emissions')  
plt.show()
```



```
[14]: from sklearn import linear_model
Regression = linear_model.LinearRegression()
train_x = np.asanyarray(train[['ENGINE_SIZE']])
train_y = np.asanyarray(train[['CO2EMISSIONS']])
Regression.fit(train_x, train_y)
```


```
[14]: LinearRegression()
```

```
[15]: print('coefficient', Regression.coef_)
print('intercept', Regression.intercept_)
```

```
coefficient [[39.02456543]]
intercept [125.76253682]
```

```
[16]: plt.scatter(cdf.ENGINE_SIZE, cdf.CO2EMISSIONS)
plt.plot(train_x, train_x * Regression.coef_ + Regression.intercept_, '-r')
```

```
[16]: [<matplotlib.lines.Line2D at 0x17c59c7f880>]
```



```
[17]: from sklearn.metrics import r2_score
```

```
[18]: test_x = np.asanyarray(test[['ENGINE_SIZE']])
test_y = np.asanyarray(test[['CO2_EMISSIONS']])
test_y_hat = Regression.predict(test_x)

print("Mean absolute error: %.2f" % np.mean(np.absolute(test_y_hat - test_y)))
print("Residual sum of squares (MSE): %.2f" % np.mean((test_y_hat - test_y) ** 2))
print("R2-score: %.2f" % r2_score(test_y_hat, test_y))
```

Mean absolute error: 22.09

Residual sum of squares (MSE): 811.08


R2-score: 0.74

```
[19]: #diff between linear and non-linear regression
```

```
[20]: x = np.arange(-5.0, 5.0, 0.1)


y = 2*(x) + 3
y_noise = 2 * np.random.normal(size=x.size)
ydata = y + y_noise
#plt.figure(figsize=(8,6))
plt.plot(x, ydata, 'bo')
plt.plot(x, y, 'r')
plt.ylabel('Dependent Variable')
```

```
plt.xlabel('Independent Variable')
plt.show()
```




```
[21]: #cubic
x = np.arange(-5.0, 5.0, 0.1)

y = 1*(x**3) + 1*(x**2) + 1*x + 3
y_noise = 20 * np.random.normal(size=x.size)
ydata = y + y_noise
plt.plot(x, ydata, 'bo')
plt.plot(x, y, 'r')
plt.ylabel('Dependent Variable')
plt.xlabel('Independent Variable')
plt.show()
```





```
[22]: #quadratic
x = np.arange(-5.0, 5.0, 0.1)

y = np.power(x,2)
y_noise = 2 * np.random.normal(size=x.size)
ydata = y + y_noise
plt.plot(x, ydata, 'bo')
plt.plot(x,y, 'r')
plt.ylabel('Dependent Variable')
plt.xlabel('Independent Variable')
plt.show()
```



```
[23]: #exponential  
X = np.arange(-5.0, 5.0, 0.1)  
  
Y= np.exp(X)  
  
plt.plot(X,Y)  
plt.ylabel('Dependent Variable')  
plt.xlabel('Independent Variable')  
plt.show()
```




```
[24]: #log
X = np.arange(-5.0, 5.0, 0.1)


Y = np.log(X)

plt.plot(X,Y)
plt.ylabel('Dependent Variable')
plt.xlabel('Independent Variable')
plt.show()
```

```
<ipython-input-24-cdc1416ded40>:4: RuntimeWarning: invalid value encountered in
log
  Y = np.log(X)
```



```
[25]: #sigmoidal/logistic  
X = np.arange(-5.0, 5.0, 0.1)  
  
Y = 1-4/(1+np.power(3, X-2))  
  
plt.plot(X,Y)  
plt.ylabel('Dependent Variable')  
plt.xlabel('Independent Variable')  
plt.show()
```



```
[26]: #non-linear regression
```


```
[27]: df1 = pd.read_csv('China_gdp.csv')
```

```
[28]: df1.head()
```


```
[28]:
```

	Year	Value
0	1960	5.918412e+10
1	1961	4.955705e+10
2	1962	4.668518e+10
3	1963	5.009730e+10
4	1964	5.906225e+10

```
[29]: plt.plot(df1.Year, df1.Value)
plt.xlabel('year')
plt.ylabel('value')
plt.show()
```



```
[30]: plt.plot(df1['Year'].values, df1['Value'].values, 'ro')
plt.xlabel('year')
plt.ylabel('value')
plt.show()
```




```
[31]: def sigmoid(x, Beta_1, Beta_2):
        y = 1 / (1 + np.exp(-Beta_1*(x-Beta_2)))
        return y
```

```
[32]: x_data = df1['Year']
      y_data = df1['Value']
```

```
[33]: beta_1 = 0.10  
beta_2 = 1990.0  
  
#logistic function  
Y_pred = sigmoid(x_data, beta_1 , beta_2)  
  
#plot initial prediction against datapoints  
plt.plot(x_data, Y_pred*150000000000000.)  
plt.plot(x_data, y_data, 'ro')
```

```
[33]: [matplotlib.lines.Line2D at 0x17c59edbbb0]
```




```
[34]: # Lets normalize our data
xdata = x_data/max(x_data)
ydata = y_data/max(y_data)
```

```
[35]: #clearly parameter doesnt fit our model. So, we'll use curve_fit
```

```
[36]: from scipy.optimize import curve_fit
popt, pcov = curve_fit(sigmoid, xdata, ydata)
#print the final parameters
print(" beta_1 = %f, beta_2 = %f" % (popt[0], popt[1]))
```

```
beta_1 = 690.451711, beta_2 = 0.997207
```

```
[37]: x = np.linspace(1960, 2015, 55)
x = x/max(x)
plt.figure(figsize=(8,5))
y = sigmoid(x, *popt)
plt.plot(xdata, ydata, 'ro', label='data')
plt.plot(x,y, linewidth=3.0, label='fit')
plt.legend(loc='best')
plt.ylabel('GDP')
plt.xlabel('Year')
plt.show()
```

```
[38]: #finding accuracy

# split data into train/test
msk = np.random.rand(len(df1)) < 0.8
train_x = xdata[msk]
test_x = xdata[~msk]
train_y = ydata[msk]
test_y = ydata[~msk]

# build the model using train set
popt, pcov = curve_fit(sigmoid, train_x, train_y)

# predict using test set
y_hat = sigmoid(test_x, *popt)

# evaluation
print("Mean absolute error: %.2f" % np.mean(np.absolute(y_hat - test_y)))
print("Residual sum of squares (MSE): %.2f" % np.mean((y_hat - test_y) ** 2))
from sklearn.metrics import r2_score
print("R2-score: %.2f" % r2_score(y_hat , test_y) )
```

```
Mean absolute error: 0.04
Residual sum of squares (MSE): 0.00
R2-score: 0.96
```

```
[39]: #Classification
```

```
[40]: #KNN
```

```
[41]: import itertools
from matplotlib.ticker import NullFormatter
import matplotlib.ticker as ticker
from sklearn import preprocessing
```

```
[42]: df2 = pd.read_csv('teleCust1000t.csv')
df2.head()
```

```
[42]:
```

	region	tenure	age	marital	address	income	ed	employ	retire	gender	\
0	2	13	44	1	9	64.0	4	5	0.0	0	
1	3	11	33	1	7	136.0	5	5	0.0	0	
2	3	68	52	1	24	116.0	1	29	0.0	1	
3	2	33	33	0	12	33.0	2	0	0.0	1	
4	2	23	30	1	9	30.0	1	2	0.0	0	

	reside	custcat
0	2	1
1	6	4
2	2	3
3	1	1
4	4	3

```
[43]: df2['custcat'].value_counts()
```


```
[43]:
```

3	281
1	266
4	236
2	217

Name: custcat, dtype: int64


```
[44]: df2.hist(column='income')
```

```
[44]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x0000017C5A38C640>]],
      dtype=object)
```



```
[45]: df2.hist(column='income', bins=50)
```

```
[45]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x0000017C5A412310>]],  
      dtype=object)
```



```
[46]: df2.columns
```

```
[46]: Index(['region', 'tenure', 'age', 'marital', 'address', 'income', 'ed',  
        'employ', 'retire', 'gender', 'reside', 'custcat'],  
        dtype='object')
```

```
[47]: #converting pandas data frame into numpys array  
X = df2[['region', 'tenure', 'age', 'marital', 'address', 'income', 'ed',  
        ↪ 'employ', 'retire', 'gender', 'reside']] .values #.astype(float)  
X[0:5]
```

```
[47]: array([[ 2., 13., 44., 1., 9., 64., 4., 5., 0., 0., 2.],  
        [ 3., 11., 33., 1., 7., 136., 5., 5., 0., 0., 6.],  
        [ 3., 68., 52., 1., 24., 116., 1., 29., 0., 1., 2.],  
        [ 2., 33., 33., 0., 12., 33., 2., 0., 0., 1., 1.],  
        [ 2., 23., 30., 1., 9., 30., 1., 2., 0., 0., 4.]])
```

```
[48]: y = df2['custcat'].values  
y[0:5]
```

```
[48]: array([1, 4, 3, 1, 3], dtype=int64)
```

```
[49]: #Data Standardization give data zero mean and unit variance, it is good  
        ↪ practice,  
        #especially for algorithms such as KNN which is based on distance of cases:
```

```
[50]: X = preprocessing.StandardScaler().fit(X).transform(X.astype(float))  
X[0:5]
```

```
[50]: array([[ -0.02696767, -1.055125 , 0.18450456, 1.0100505 , -0.25303431,  
        -0.12650641, 1.0877526 , -0.5941226 , -0.22207644, -1.03459817,  
        -0.23065004],  
        [ 1.19883553, -1.14880563, -0.69181243, 1.0100505 , -0.4514148 ,  
        0.54644972, 1.9062271 , -0.5941226 , -0.22207644, -1.03459817,  
        2.55666158],  
        [ 1.19883553, 1.52109247, 0.82182601, 1.0100505 , 1.23481934,  
        0.35951747, -1.36767088, 1.78752803, -0.22207644, 0.96655883,  
        -0.23065004],  
        [ -0.02696767, -0.11831864, -0.69181243, -0.9900495 , 0.04453642,  
        -0.41625141, -0.54919639, -1.09029981, -0.22207644, 0.96655883,  
        -0.92747794],  
        [ -0.02696767, -0.58672182, -0.93080797, 1.0100505 , -0.25303431,  
        -0.44429125, -1.36767088, -0.89182893, -0.22207644, -1.03459817,  
        1.16300577]])
```

```
[51]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split( X, y, test_size=0.2,
↳random_state=4)
print ('Train set:', X_train.shape, y_train.shape)
print ('Test set:', X_test.shape, y_test.shape)
```

```
Train set: (800, 11) (800,)
Test set: (200, 11) (200,)
```

```
[52]: from sklearn.neighbors import KNeighborsClassifier
```

```
[53]: neigh = KNeighborsClassifier(n_neighbors=4)
```

```
[54]: neigh.fit(X_train, y_train)
```

```
[54]: KNeighborsClassifier(n_neighbors=4)
```

```
[55]: yhat = neigh.predict(X_test)
```

```
[56]: from sklearn.metrics import accuracy_score
```

```
[57]: accuracy_score(yhat, y_test)
```

```
[57]: 0.32
```

```
[58]: from sklearn import metrics
print("Train set Accuracy: ", metrics.accuracy_score(y_train, neigh.
↳predict(X_train)))
print("Test set Accuracy: ", metrics.accuracy_score(y_test, yhat))
```

```
Train set Accuracy: 0.5475
Test set Accuracy: 0.32
```

```
[59]: #to see which k is the best
Ks = 10
mean_acc = np.zeros((Ks-1))
std_acc = np.zeros((Ks-1))
ConfustionMx = [];
for n in range(1,Ks):


    #Train Model and Predict
    neigh = KNeighborsClassifier(n_neighbors = n).fit(X_train,y_train)
    yhat=neigh.predict(X_test)
    mean_acc[n-1] = metrics.accuracy_score(y_test, yhat)

    std_acc[n-1]=np.std(yhat==y_test)/np.sqrt(yhat.shape[0])
```

```
mean_acc
```

```
[59]: array([0.3 , 0.29 , 0.315, 0.32 , 0.315, 0.31 , 0.335, 0.325, 0.34 ])
```

```
[60]: plt.plot(range(1,Ks),mean_acc,'g')
plt.fill_between(range(1,Ks),mean_acc - 1 * std_acc,mean_acc + 1 * std_acc,
    ↪alpha=0.10)
plt.legend(('Accuracy ', '+/- 3xstd'))
plt.ylabel('Accuracy ')
plt.xlabel('Number of Neighbors (K)')
plt.tight_layout()
plt.show()
```



```
[61]: #The best accuracy was with k=9
```

```
[62]: #DECISION TREE
```

```
[99]: from sklearn.tree import DecisionTreeClassifier
```

```
[65]: df3 = pd.read_csv('drug200.csv')
df3.head()
```

```
[65]:
```

	Age	Sex	BP	Cholesterol	Na_to_K	Drug
0	23	F	HIGH	HIGH	25.355	drugY
1	47	M	LOW	HIGH	13.093	drugC
2	47	M	LOW	HIGH	10.114	drugC
3	28	F	NORMAL	HIGH	7.798	drugX
4	61	F	LOW	HIGH	18.043	drugY

```
[82]: x = df3[['Age', 'Sex', 'BP', 'Cholesterol', 'Na_to_K']].values
```

```
[83]: #cleaning categorical data
```

```
[84]: from sklearn import preprocessing

le_sex = preprocessing.LabelEncoder()
le_sex.fit(['F', 'M'])
x[:, 1] = le_sex.transform(x[:, 1])
```

```
[86]: le_bp = preprocessing.LabelEncoder()
le_bp.fit(['HIGH', 'LOW', 'NORMAL'])
x[:, 2] = le_bp.transform(x[:, 2])
```

```
[87]: le_cholesterol = preprocessing.LabelEncoder()
le_cholesterol.fit(['NORMAL', 'HIGH'])
x[:, 3] = le_cholesterol.transform(x[:, 3])
```

```
[92]: x[0:5]
```

```
[92]: array([[23, 0, 0, 0, 25.355],
          [47, 1, 1, 0, 13.093],
          [47, 1, 1, 0, 10.113999999999999],
          [28, 0, 2, 0, 7.797999999999999],
          [61, 0, 1, 0, 18.043]], dtype=object)
```

```
[93]: y = df3['Drug']
```

```
[94]:
```

```
[94]: 0      drugY
      1      drugC
      2      drugC
      3      drugX
      4      drugY
      ...
     195     drugC
     196     drugC
     197     drugX
     198     drugX
     199     drugX
```

Name: Drug, Length: 200, dtype: object

```
[97]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
↳ random_state=3)
```

```
[117]: tree1 = DecisionTreeClassifier(criterion='entropy', max_depth=4)
tree1
```

```
[117]: DecisionTreeClassifier(criterion='entropy', max_depth=4)
```

```
[118]: tree1.fit(x_train, y_train)
```

```
[118]: DecisionTreeClassifier(criterion='entropy', max_depth=4)
```

```
[119]: prediction = tree1.predict(x_test)
```

```
[120]: Accuracy = accuracy_score(prediction, y_test)
Accuracy
```

```
[120]: 1.0
```

```
[121]: import io
import pydotplus
import matplotlib.image as mpimg
from sklearn import tree
%matplotlib inline
```

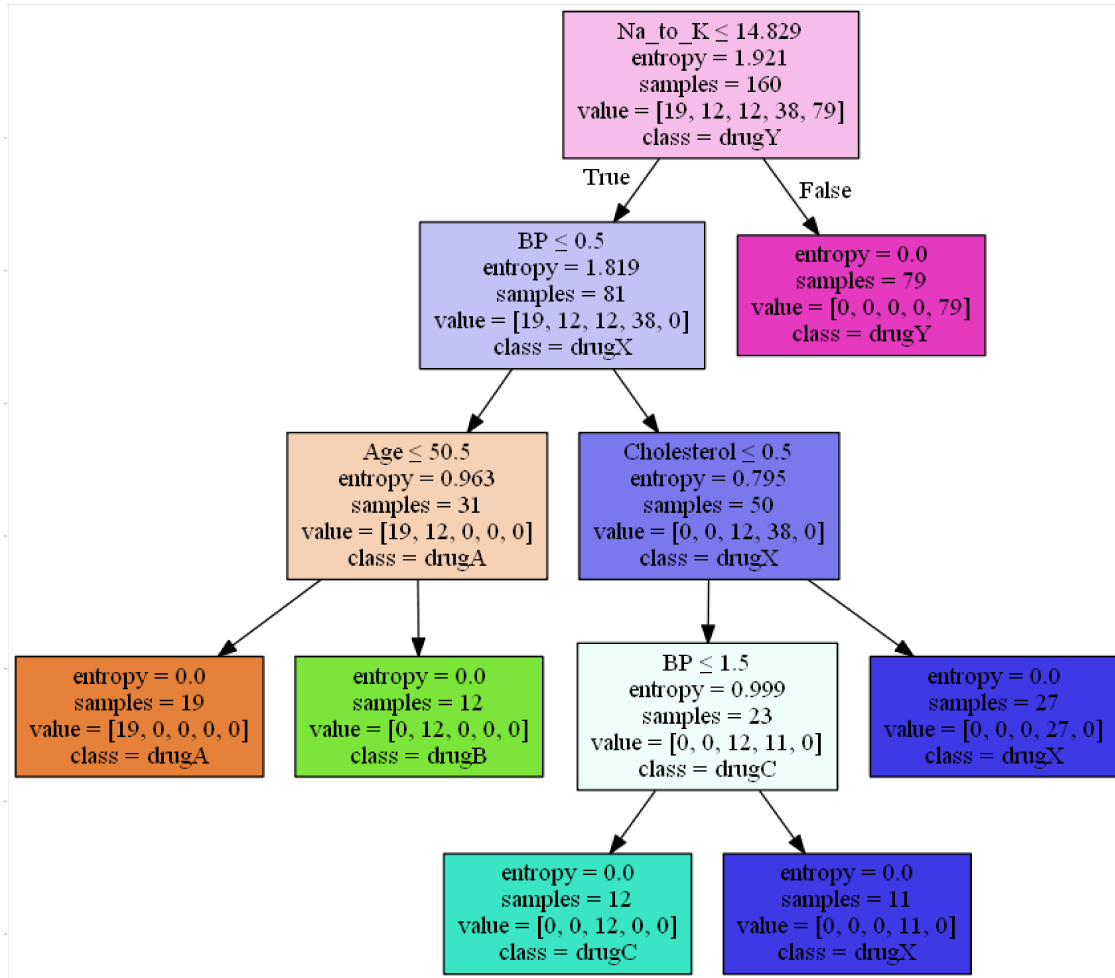
```
[ ]: estimator = model.estimators_[5]
export_graphviz(estimator_limited,
                out_file='tree.dot',
                feature_names = iris.feature_names,
                class_names = iris.target_names,
                rounded = True, proportion = False,
                precision = 2, filled = True)
call(['dot', '-Tpng', 'tree.dot', '-o', 'tree.png', '-Gdpi=600'])
from IPython.display import Image
Image(filename = 'tree.png')
```

```
[123]: dot_data = io.StringIO()
filename = "drugtree.png"
featureNames = df3.columns[0:5]
targetNames = df3["Drug"].unique().tolist()
out=tree.export_graphviz(tree1,feature_names=featureNames, out_file=dot_data,
↳ class_names= np.unique(y_train), filled=True,
↳ special_characters=True,rotate=False)
graph = pydotplus.graph_from_dot_data(dot_data.getvalue())
graph.write_png(filename)
img = mpimg.imread(filename)
```



```
plt.figure(figsize=(100, 200))
plt.imshow(img, interpolation='nearest')
```

[123]: <matplotlib.image.AxesImage at 0x17c5a8d9550>



[124]: *#LOGISTIC REGRESSION*

```
[125]: import pandas as pd
import pylab as pl
import numpy as np
import scipy.optimize as opt
from sklearn import preprocessing
%matplotlib inline
import matplotlib.pyplot as plt
```

[126]: df4 = pd.read_csv('ChurnData.csv')

```
[127]: df4.head()
```

```
[127]:
```

	tenure	age	address	income	ed	employ	equip	callcard	wireless	\
0	11.0	33.0	7.0	136.0	5.0	5.0	0.0	1.0	1.0	
1	33.0	33.0	12.0	33.0	2.0	0.0	0.0	0.0	0.0	
2	23.0	30.0	9.0	30.0	1.0	2.0	0.0	0.0	0.0	
3	38.0	35.0	5.0	76.0	2.0	10.0	1.0	1.0	1.0	
4	7.0	35.0	14.0	80.0	2.0	15.0	0.0	1.0	0.0	

	longmon	...	pager	internet	callwait	confer	ebill	loglong	logtoll	\
0	4.40	...	1.0	0.0	1.0	1.0	0.0	1.482	3.033	
1	9.45	...	0.0	0.0	0.0	0.0	0.0	2.246	3.240	
2	6.30	...	0.0	0.0	0.0	1.0	0.0	1.841	3.240	
3	6.05	...	1.0	1.0	1.0	1.0	1.0	1.800	3.807	
4	7.10	...	0.0	0.0	1.0	1.0	0.0	1.960	3.091	

	lninc	custcat	churn
0	4.913	4.0	1.0
1	3.497	1.0	1.0
2	3.401	3.0	0.0
3	4.331	4.0	0.0
4	4.382	3.0	0.0

[5 rows x 28 columns]

```
[131]: df4 = df4[['tenure', 'age', 'address', 'income', 'ed', 'employ', 'equip', 'callcard', 'wireless', 'churn']]
df4['churn'] = df4['churn'].astype('int')
df4.head()
```

```
[131]:
```

	tenure	age	address	income	ed	employ	equip	callcard	wireless	\
0	11.0	33.0	7.0	136.0	5.0	5.0	0.0	1.0	1.0	
1	33.0	33.0	12.0	33.0	2.0	0.0	0.0	0.0	0.0	
2	23.0	30.0	9.0	30.0	1.0	2.0	0.0	0.0	0.0	
3	38.0	35.0	5.0	76.0	2.0	10.0	1.0	1.0	1.0	
4	7.0	35.0	14.0	80.0	2.0	15.0	0.0	1.0	0.0	

	churn
0	1
1	1
2	0
3	0
4	0

```
[134]: X = df4[['tenure', 'age', 'address', 'income', 'ed', 'employ', 'equip']].values
X[0:5]
```

```
[134]: array([[ 11.,  33.,   7., 136.,   5.,   5.,   0.],
              [ 33.,  33.,  12.,  33.,   2.,   0.,   0.],
              [ 23.,  30.,   9.,  30.,   1.,   2.,   0.],
              [ 38.,  35.,   5.,  76.,   2.,  10.,   1.],
              [  7.,  35.,  14.,  80.,   2.,  15.,   0.]])
```

```
[136]: X = np.asarray(df4[['tenure', 'age', 'address', 'income', 'ed', 'employ', 'equip']])
X[0:5]
```

```
[136]: array([[ 11.,  33.,   7., 136.,   5.,   5.,   0.],
              [ 33.,  33.,  12.,  33.,   2.,   0.,   0.],
              [ 23.,  30.,   9.,  30.,   1.,   2.,   0.],
              [ 38.,  35.,   5.,  76.,   2.,  10.,   1.],
              [  7.,  35.,  14.,  80.,   2.,  15.,   0.]])
```

```
[139]: y = np.asarray(df4['churn'])
y[0:5]
```

```
[139]: array([1, 1, 0, 0, 0])
```

```
[141]: X = preprocessing.StandardScaler().fit(X).transform(X)
X[0:5]
```

```
[141]: array([[ -1.13518441, -0.62595491, -0.4588971 ,  0.4751423 ,  1.6961288 ,
               -0.58477841, -0.85972695],
              [-0.11604313, -0.62595491,  0.03454064, -0.32886061, -0.6433592 ,
               -1.14437497, -0.85972695],
              [-0.57928917, -0.85594447, -0.261522  , -0.35227817, -1.42318853,
               -0.92053635, -0.85972695],
              [ 0.11557989, -0.47262854, -0.65627219,  0.00679109, -0.6433592 ,
               -0.02518185,  1.16316  ],
              [-1.32048283, -0.47262854,  0.23191574,  0.03801451, -0.6433592 ,
               0.53441472, -0.85972695]])
```

```
[142]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
               random_state=4)
```

```
[143]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix
```

```
[144]: LR = LogisticRegression()
```

```
[145]: LR.fit(X_train, y_train)
```

```
[145]: LogisticRegression()
```

```
[147]: Yhat = LR.predict(X_test)
```

```
[163]: y_test.shape
```

```
[163]: (40,)
```

```
[148]: #predict_proba returns estimates for all classes, ordered by the label of  
       ↪ classes.  
       # So, the first column is the probability of class 1,  $P(Y=1/X)$ , and second  
       ↪ column is probability of class 0,  $P(Y=0/X)$ 
```

```
[152]: yhat_prob = LR.predict_proba(X_test)  
       yhat_prob
```

```
[152]: array([[0.74658429, 0.25341571],  
             [0.92677899, 0.07322101],  
             [0.83445726, 0.16554274],  
             [0.94596742, 0.05403258],  
             [0.84351139, 0.15648861],  
             [0.71452329, 0.28547671],  
             [0.77085785, 0.22914215],  
             [0.90956492, 0.09043508],  
             [0.26142925, 0.73857075],  
             [0.94907369, 0.05092631],  
             [0.84772942, 0.15227058],  
             [0.89315103, 0.10684897],  
             [0.57506834, 0.42493166],  
             [0.32555873, 0.67444127],  
             [0.91995311, 0.08004689],  
             [0.633071  , 0.366929  ],  
             [0.6297197 , 0.3702803 ],  
             [0.71293143, 0.28706857],  
             [0.64068923, 0.35931077],  
             [0.7794542 , 0.2205458 ],  
             [0.91593448, 0.08406552],  
             [0.64123809, 0.35876191],  
             [0.96435248, 0.03564752],  
             [0.55216187, 0.44783813],  
             [0.62291087, 0.37708913],  
             [0.97603043, 0.02396957],  
             [0.6014112 , 0.3985888 ],  
             [0.68062074, 0.31937926],  
             [0.71779212, 0.28220788],  
             [0.9820836 , 0.0179164 ],  
             [0.96445529, 0.03554471],  
             [0.765139  , 0.234861  ],  
             [0.29422866, 0.70577134],  
             [0.96537173, 0.03462827],  
             [0.93653126, 0.06346874],
```

```
[0.88756299, 0.11243701],
[0.22268043, 0.77731957],
[0.70018568, 0.29981432],
[0.93940692, 0.06059308],
[0.72116371, 0.27883629]])
```

```
[1]: #SVM
```

```
[2]: import pandas as pd
import pylab as pl
import numpy as np
import scipy.optimize as opt
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
%matplotlib inline
```

```
[5]: df = pd.read_csv('cell_samples.csv')
df.tail()
```

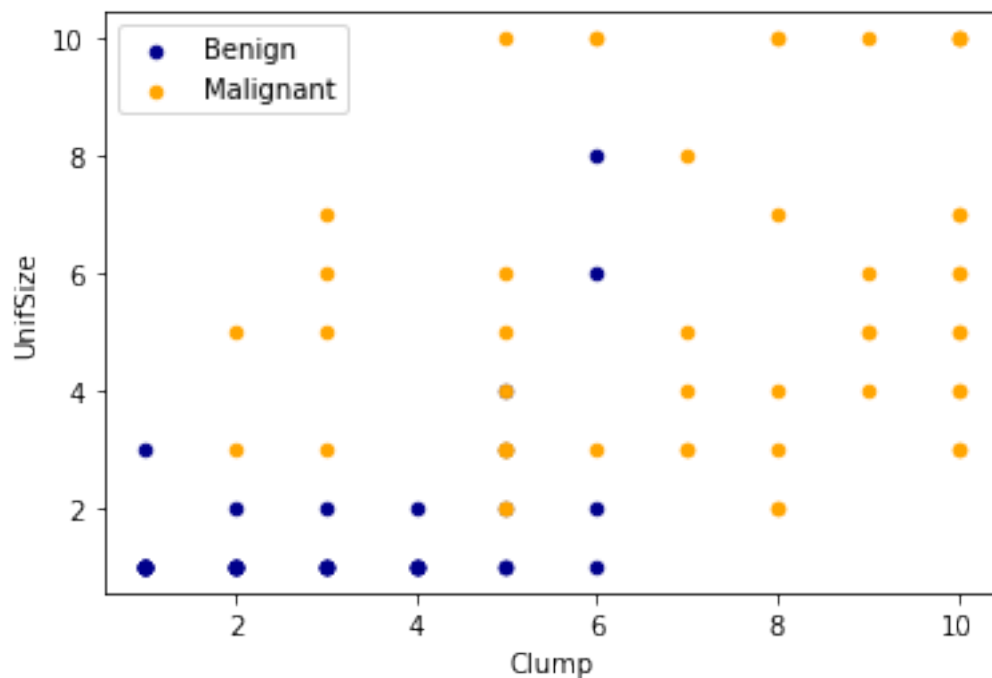
```
[5]:
```

	ID	Clump	UnifSize	UnifShape	MargAdh	SingEpiSize	BareNuc	\
694	776715	3	1	1	1	3	2	
695	841769	2	1	1	1	2	1	
696	888820	5	10	10	3	7	3	
697	897471	4	8	6	4	3	4	
698	897471	4	8	8	5	4	5	

	BlandChrom	NormNucl	Mit	Class
694	1	1	1	2
695	1	1	1	2
696	8	10	2	4
697	10	6	1	4
698	10	4	1	4

```
[6]: #The attribute class has 2 and 4 where 2 means benign and 4 means malignant
      ↪type of cancer
```

```
[16]: ax = df[df['Class'] == 2][0:50].plot(kind='scatter', x='Clump', y='UnifSize',
      ↪color='DarkBlue', label='Benign');
df[df['Class'] == 4][0:50].plot(kind='scatter', x='Clump', y='UnifSize',
      ↪color='Orange', label='Malignant', ax=ax);
plt.show()
```



```
[17]: df.dtypes
```

```
[17]: ID                int64
      Clump             int64
      UnifSize          int64
      UnifShape          int64
      MargAdh           int64
      SingEpiSize       int64
      BareNuc           object
      BlandChrom         int64
      NormNucl          int64
      Mit               int64
      Class             int64
      dtype: object
```

```
[20]: df = df[pd.to_numeric(df['BareNuc'], errors='coerce').notnull()]
      df['BareNuc'] = df['BareNuc'].astype('int')
      df.dtypes
```

```
[20]: ID                int64
      Clump             int64
      UnifSize          int64
      UnifShape          int64
      MargAdh           int64
      SingEpiSize       int64
```

```
BareNuc      int32
BlandChrom   int64
NormNucl     int64
Mit          int64
Class        int64
dtype: object
```

```
[29]: feature_df = df.drop(columns='Class')
      X = np.asarray(feature_df)
```

[31] : X

```
[31]: array([[1000025,      5,      1, ...,      3,      1,      1],
             [1002945,      5,      4, ...,      3,      2,      1],
             [1015425,      3,      1, ...,      3,      1,      1],
             ...,
             [ 888820,      5,     10, ...,      8,     10,      2],
             [ 897471,      4,      8, ...,     10,      6,      1],
             [ 897471,      4,      8, ...,     10,      4,      1]],
          dtype=int64)
```

```
[32]: y = df['Class'].values
```

[34] : y

[illegible]

```

2, 2, 2, 2, 2, 2, 2, 4, 2, 2, 2, 2, 2, 2, 2, 4, 4, 2, 2, 2, 4, 2,
2, 4, 4, 2, 2, 2, 2, 2, 2, 4, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
2, 2, 2, 4, 2, 2, 4, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,
4, 2, 2, 4, 4, 4, 4, 2, 2, 4, 2, 2, 2, 2, 2, 2, 4, 4, 2, 2, 2, 4,
2, 4, 2, 4, 4, 4, 2, 4, 2, 2, 2, 2, 2, 2, 2, 2, 4, 4, 4, 2, 2, 4,
2, 4, 4, 4, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 4, 2, 2, 2, 2, 2,
2, 4, 2, 2, 4, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 4, 2, 2, 2, 2, 2,
2, 2, 2, 2, 4, 2, 2, 2, 2, 2, 2, 2, 2, 2, 4, 4, 4, 2, 2, 2, 2, 2,
2, 2, 2, 2, 4, 4, 2, 2, 2, 2, 2, 2, 2, 2, 4, 2, 2, 2, 2, 4, 4,
4], dtype=int64)

```

```
[35]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=4)
```

```
[36]: from sklearn import svm
```

```
[38]: model = svm.SVC(kernel='rbf')
```

```
[39]: model.fit(X_train, y_train)
```

```
[39]: SVC()
```

```
[40]: yhat = model.predict(X_test)
```

```
[53]: from sklearn.metrics import accuracy_score
score = accuracy_score(y_test, yhat)
score
```

```
[53]: 0.656934306569343
```

```
[50]: from sklearn.metrics import f1_score
f1_score(y_test, yhat, average='weighted')
```

```
[50]: 0.5209170712884659
```

```
[54]: model = svm.SVC(kernel='linear')
model.fit(X_train, y_train)
yhat = model.predict(X_test)
```

```
[55]: from sklearn.metrics import accuracy_score
score = accuracy_score(y_test, yhat)
score
```

```
[55]: 0.6861313868613139
```

```
[1]: #K-means using random generated datasets
```



```
[6]: import numpy as np
import random
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.datasets.samples_generator import make_blobs
%matplotlib inline
```


C:\Users\Dell\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:143:
FutureWarning: The sklearn.datasets.samples_generator module is deprecated in
version 0.22 and will be removed in version 0.24. The corresponding classes /
functions should instead be imported from sklearn.datasets. Anything that cannot
be imported from sklearn.datasets is now part of the private API.
warnings.warn(message, FutureWarning)

```
[9]: np.random.seed(0)
```

```
[12]: X,y =make_blobs(n_samples = 5000, centers = [[4,4], [2,-1], [2,-3], [1,1]],  
→cluster_std=0.9)
```

```
[20]: plt.scatter(X[:,0], X[:,1], marker='o')
```

```
[20]: <matplotlib.collections.PathCollection at 0x2bec1c87250>
```



```
[23]: k_means = KMeans(init='k-means++', n_clusters=4, n_init=12)
```

```
[24]: k_means.fit(X)
```

```
[24]: KMeans(n_clusters=4, n_init=12)
```

```
[25]: k_means_labels = k_means.labels_  
k_means_labels
```

```
[25]: array([0, 1, 1, ..., 2, 0, 3])
```

```
[26]: k_means_cluster_center = k_means.cluster_centers_  
k_means_cluster_center
```

```
[26]: array([[ 1.99386362, -0.98210637],  
            [ 2.00767516, -3.16352369],  
            [ 3.97423467,  3.98310553],  
            [ 0.90429725,  1.07253251]])
```

```
[28]: # Initialize the plot with the specified dimensions.  
fig = plt.figure(figsize=(6, 4))  
  
# Colors uses a color map, which will produce an array of colors based on  
# the number of labels there are. We use set(k_means_labels) to get the  
# unique labels.  
colors = plt.cm.Spectral(np.linspace(0, 1, len(set(k_means_labels))))  
  
# Create a plot  
ax = fig.add_subplot(1, 1, 1)  
  
# For loop that plots the data points and centroids.  
# k will range from 0-3, which will match the possible clusters that each  
# data point is in.  
for k, col in zip(range(len([[4,4], [-2, -1], [2, -3], [1, 1]])), colors):  
  
    # Create a list of all data points, where the data points that are  
    # in the cluster (ex. cluster 0) are labeled as true, else they are  
    # labeled as false.  
    my_members = (k_means_labels == k)  
  
    # Define the centroid, or cluster center.  
    cluster_center = k_means_cluster_center[k]  
  
    # Plots the datapoints with color col.  
    ax.plot(X[my_members, 0], X[my_members, 1], 'w', markerfacecolor=col,  
↪marker='.')  
  
    # Plots the centroids with specified color, but with a darker outline  
    ax.plot(cluster_center[0], cluster_center[1], 'o', markerfacecolor=col, ↪  
↪markeredgecolor='k', markersize=6)
```

```


# Title of the plot
ax.set_title('KMeans')

# Remove x-axis ticks
ax.set_xticks(())

# Remove y-axis ticks
ax.set_yticks(())

# Show the plot
plt.show()

```



```
[29]: #k-means using dataset
```

```

[30]: import pandas as pd
      cust_df = pd.read_csv("Cust_Segmentation.csv")
      cust_df.head()

```

```

[30]:   Customer Id  Age  Edu  Years Employed  Income  Card Debt  Other Debt  \
0         1    41    2         6         19      0.124      1.073
1         2    47    1        26        100      4.582      8.218
2         3    33    2        10         57      6.111      5.802
3         4    29    2         4         19      0.681      0.516
4         5    47    1        31        253      9.308      8.908

```

	Defaulted	Address	DebtIncomeRatio
0	0.0	NBA001	6.3
1	0.0	NBA021	12.8
2	1.0	NBA013	20.9
3	0.0	NBA009	6.3
4	0.0	NBA008	7.2

```
[32]: cust_df.drop('Address', axis=1, inplace=True)
```

```
[33]: cust_df
```

```
[33]:
```

	Customer	Id	Age	Edu	Years	Employed	Income	Card Debt	Other Debt	\
0		1	41	2		6	19	0.124	1.073	
1		2	47	1		26	100	4.582	8.218	
2		3	33	2		10	57	6.111	5.802	
3		4	29	2		4	19	0.681	0.516	
4		5	47	1		31	253	9.308	8.908	
..			
845		846	27	1		5	26	0.548	1.220	
846		847	28	2		7	34	0.359	2.021	
847		848	25	4		0	18	2.802	3.210	
848		849	32	1		12	28	0.116	0.696	
849		850	52	1		16	64	1.866	3.638	

	Defaulted	DebtIncomeRatio
0	0.0	6.3
1	0.0	12.8
2	1.0	20.9
3	0.0	6.3
4	0.0	7.2
..
845	NaN	6.8
846	0.0	7.0
847	1.0	33.4
848	0.0	2.9
849	0.0	8.6

[850 rows x 9 columns]

```
[35]: from sklearn.preprocessing import StandardScaler
X = cust_df.values[:,1:]
X = np.nan_to_num(X)
Clus_dataSet = StandardScaler().fit_transform(X)
Clus_dataSet
```

```
[35]: array([[ 0.74291541,  0.31212243, -0.37878978, ..., -0.59048916,
           -0.52379654, -0.57652509],
```

```
[ 1.48949049, -0.76634938,  2.5737211 , ...,  1.51296181,
 -0.52379654,  0.39138677],
 [-0.25251804,  0.31212243,  0.2117124 , ...,  0.80170393,
  1.90913822,  1.59755385],
 ...,
 [-1.24795149,  2.46906604, -1.26454304, ...,  0.03863257,
  1.90913822,  3.45892281],
 [-0.37694723, -0.76634938,  0.50696349, ..., -0.70147601,
 -0.52379654, -1.08281745],
 [ 2.1116364 , -0.76634938,  1.09746566, ...,  0.16463355,
 -0.52379654, -0.2340332 ]])
```

```
[36]: clusterNum = 3
k_means = KMeans(init = "k-means++", n_clusters = clusterNum, n_init = 12)
k_means.fit(X)
labels = k_means.labels_
print(labels)
```

```
[1 0 1 1 2 0 1 0 1 0 0 1 1 1 1 1 1 1 0 1 1 1 1 0 0 0 1 1 0 1 0 1 1 1 1 1 1
 1 1 0 1 0 1 2 1 0 1 1 1 0 0 1 1 0 0 1 1 1 0 1 0 1 0 0 1 1 0 1 1 1 0 0 0 1
 1 1 1 1 0 1 0 0 2 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 0 0 1 1 1 1 1 1 0 1
 1 1 1 1 1 1 1 0 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 0 1 0 1
 1 1 1 1 1 1 0 1 0 0 1 0 1 1 0 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 0 1 1 1 0 1
 1 1 1 1 0 1 1 0 1 0 1 1 0 2 1 0 1 1 1 1 1 1 2 0 1 1 1 1 0 1 1 0 0 1 0 1 0
 1 1 1 1 0 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 2 0 1 1 1 1 1 1 1 0 1 1 1 1
 1 1 0 1 1 0 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 0 1 0 1 0 1 0 0 1 1 1 1 1
 1 1 1 0 0 0 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 0 1 1 1 1 1 0 1 0 0 1
 1 1 1 1 0 1 1 1 1 1 1 0 1 1 0 1 1 0 1 1 1 1 1 0 1 1 1 2 1 1 1 0 1 0 0 0 1
 1 1 0 1 1 1 1 1 1 1 1 1 1 1 0 1 0 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1
 1 0 1 1 0 1 1 1 1 0 1 1 1 1 0 1 1 0 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 0 1 1 1 2
 1 1 1 1 1 1 0 1 1 1 2 1 1 1 1 0 1 2 1 1 1 1 0 1 0 0 0 1 1 0 0 1 1 1 1 1 1
 1 0 1 1 1 1 0 1 1 1 0 1 0 1 1 1 0 1 1 1 1 0 0 1 1 1 1 0 1 1 1 1 0 1 1 1 1
 1 0 0 1 1 1 1 1 1 1 1 1 1 1 2 0 1 1 1 1 1 1 1 0 1 1 1 1 0 1 1 0 1 1 2 1 2 1
 1 2 1 1 1 1 1 1 1 1 1 0 1 0 1 1 2 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 0 1 0
 1 1 1 1 1 1 0 1 1 1 1 0 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 0
 0 1 1 0 1 0 1 1 0 1 0 1 1 2 1 0 1 0 1 1 1 1 1 0 0 1 1 1 1 0 1 1 1 0 0 1 1
 0 1 1 1 0 1 2 1 1 0 1 1 1 1 1 1 1 0 1 1 1 0 1 1 1 1 1 0 1 1 0 1 1 1 1 1 1
 1 1 0 1 1 0 1 0 1 0 0 1 1 1 0 1 0 1 1 1 1 1 0 1 1 1 1 0 0 1 1 0 0 1 1 1 1
 1 0 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 0 1 0 0 1 0 1 0 0 1 1 0 1 1 1 1 1 0 0
 1 1 1 1 1 1 1 0 1 1 1 1 1 1 2 0 0 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 0 1 1 1 1
 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 0]
```

```
[38]: cust_df["Clus_km"] = labels
cust_df.head(5)
```

```
[38]:   Customer Id  Age  Edu  Years Employed  Income  Card Debt  Other Debt  \
0          1    41    2           6         19      0.124      1.073
1          2    47    1          26        100      4.582      8.218
```

2	3	33	2	10	57	6.111	5.802
3	4	29	2	4	19	0.681	0.516
4	5	47	1	31	253	9.308	8.908

	Defaulted	DebtIncomeRatio	Clus_km
0	0.0	6.3	1
1	0.0	12.8	0
2	1.0	20.9	1
3	0.0	6.3	1
4	0.0	7.2	2

```
[40]: cust_df.groupby('Clus_km').mean()
```


```
[40]:
```

	Customer Id	Age	Edu	Years Employed	Income \
Clus_km					
0	403.780220	41.368132	1.961538	15.252747	84.076923
1	432.006154	32.967692	1.613846	6.389231	31.204615
2	410.166667	45.388889	2.666667	19.555556	227.166667

	Card Debt	Other Debt	Defaulted	DebtIncomeRatio
Clus_km				
0	3.114412	5.770352	0.172414	10.725824
1	1.032711	2.108345	0.284658	10.095385
2	5.678444	10.907167	0.285714	7.322222

```
[41]: area = np.pi * ( X[:, 1])**2
plt.scatter(X[:, 0], X[:, 3], s=area, c=labels.astype(np.float), alpha=0.5)
plt.xlabel('Age', fontsize=18)
plt.ylabel('Income', fontsize=16)

plt.show()
```




```
[42]: from mpl_toolkits.mplot3d import Axes3D
fig = plt.figure(1, figsize=(8, 6))
plt.clf()
ax = Axes3D(fig, rect=[0, 0, .95, 1], elev=48, azim=134)

plt.cla()
# plt.ylabel('Age', fontsize=18)
# plt.xlabel('Income', fontsize=16)
# plt.zlabel('Education', fontsize=16)
ax.set_xlabel('Education')
ax.set_ylabel('Age')
ax.set_zlabel('Income')

ax.scatter(X[:, 1], X[:, 0], X[:, 3], c= labels.astype(np.float))
```

```
[42]: <mpl_toolkits.mplot3d.art3d.Path3DCollection at 0x2bec34bc250>
```




```
[1]: #hierarchial clustering using random data
```

```
[4]: import numpy as np
import pandas as pd
from scipy import ndimage
from scipy.cluster import hierarchy
from scipy.spatial import distance_matrix
from matplotlib import pyplot as plt
from sklearn import manifold, datasets
from sklearn.cluster import AgglomerativeClustering
from sklearn.datasets.samples_generator import make_blobs
%matplotlib inline
```

```
[5]: X1,y1 = make_blobs(n_samples=50, centers=[[4,4],[-2,-1],[1,1],[10,4]],
↳cluster_std=0.9)
```

```
[6]: plt.scatter(X1[:,0], X1[:,1], marker='.'))
```

```
[6]: <matplotlib.collections.PathCollection at 0x1feeca44fd0>
```

```
[7]: agglom = AgglomerativeClustering(n_clusters = 4, linkage = 'average')
```

```
[8]: agglom.fit(X1,y1)
```

```
[8]: AgglomerativeClustering(linkage='average', n_clusters=4)
```

```
[10]: dist_matrix = distance_matrix(X1,X1)
print(dist_matrix)
```


```
[[0.          0.12247938 0.74991261 ... 0.76668181 1.06019761 0.46335653]
 [0.12247938 0.          0.62989098 ... 0.65519682 0.9547588  0.36737556]
 [0.74991261 0.62989098 0.          ... 0.15898383 0.42069497 0.36995815]
 ...
 [0.76668181 0.65519682 0.15898383 ... 0.          0.3101169  0.32386818]
 [1.06019761 0.9547588  0.42069497 ... 0.3101169  0.          0.59944437]
 [0.46335653 0.36737556 0.36995815 ... 0.32386818 0.59944437 0.          ]]
```

```
[11]: Z = hierarchy.linkage(dist_matrix, 'complete')
```

```
<ipython-input-11-3814b774a052>:1: ClusterWarning: scipy.cluster: The symmetric
non-negative hollow observation matrix looks suspiciously like an uncondensed
distance matrix
```

```
    Z = hierarchy.linkage(dist_matrix, 'complete')
```

```
[14]: dendro = hierarchy.dendrogram(Z)
```




```
[17]: Z = hierarchy.linkage(dist_matrix, 'centroid')
```

<ipython-input-17-a08c4ecedfa6>:1: ClusterWarning: scipy.cluster: The symmetric non-negative hollow observation matrix looks suspiciously like an uncondensed distance matrix

```
    Z = hierarchy.linkage(dist_matrix, 'centroid')
```

```
[18]: dendro = hierarchy.dendrogram(Z)
```



```
[19]: #hierarchial clustering on real dataset
```

```
[27]: df= pd.read_csv('cars_clus.csv')
df.head()
```

```
[27]:
```

	manufact	model	sales	resale	type	price	engine_s	horsepow	wheelbas	\
0	Acura	Integra	16.919	16.360	0.000	21.500	1.800	140.000	101.200	
1	Acura	TL	39.384	19.875	0.000	28.400	3.200	225.000	108.100	
2	Acura	CL	14.114	18.225	0.000	\$null\$	3.200	225.000	106.900	
3	Acura	RL	8.588	29.725	0.000	42.000	3.500	210.000	114.600	
4	Audi	A4	20.397	22.255	0.000	23.990	1.800	150.000	102.600	

	width	length	curb_wgt	fuel_cap	mpg	lnsales	partition
0	67.300	172.400	2.639	13.200	28.000	2.828	0.0
1	70.300	192.900	3.517	17.200	25.000	3.673	0.0
2	70.600	192.000	3.470	17.200	26.000	2.647	0.0
3	71.400	196.600	3.850	18.000	22.000	2.150	0.0
4	68.200	178.000	2.998	16.400	27.000	3.015	0.0

```
[29]: print ("Shape of dataset before cleaning: ", df.size)
df[[ 'sales', 'resale', 'type', 'price', 'engine_s',
      'horsepow', 'wheelbas', 'width', 'length', 'curb_wgt', 'fuel_cap',
      'mpg', 'lnsales']] = df[['sales', 'resale', 'type', 'price', 'engine_s',
      'horsepow', 'wheelbas', 'width', 'length', 'curb_wgt', 'fuel_cap',
      'mpg', 'lnsales']].apply(pd.to_numeric, errors='coerce')
df = df.dropna()
```

```
df = df.reset_index(drop=True)
print ("Shape of dataset after cleaning: ", df.size)
df.head(5)
```

Shape of dataset before cleaning: 2544

Shape of dataset after cleaning: 1872

```
[29]:  manufact    model    sales  resale  type  price  engine_s  horsepower \
0      Acura  Integra  16.919  16.360   0.0  21.50         1.8      140.0
1      Acura      TL   39.384  19.875   0.0  28.40         3.2      225.0
2      Acura      RL    8.588  29.725   0.0  42.00         3.5      210.0
3      Audi     A4   20.397  22.255   0.0  23.99         1.8      150.0
4      Audi     A6   18.780  23.555   0.0  33.95         2.8      200.0

      wheelbas  width  length  curb_wgt  fuel_cap  mpg  lnsales  partition
0      101.2    67.3   172.4    2.639    13.2  28.0    2.828         0.0
1      108.1    70.3   192.9    3.517    17.2  25.0    3.673         0.0
2      114.6    71.4   196.6    3.850    18.0  22.0    2.150         0.0
3      102.6    68.2   178.0    2.998    16.4  27.0    3.015         0.0
4      108.7    76.1   192.0    3.561    18.5  22.0    2.933         0.0
```

```
[30]: featureset = df[['engine_s', 'horsepow', 'wheelbas', 'width', 'length',
    ↪ 'curb_wgt', 'fuel_cap', 'mpg']]
```

```
[31]: from sklearn.preprocessing import MinMaxScaler
x = featureset.values #returns a numpy array
min_max_scaler = MinMaxScaler()
feature_mtx = min_max_scaler.fit_transform(x)
feature_mtx [0:5]
```

```
[31]: array([[0.11428571, 0.21518987, 0.18655098, 0.28143713, 0.30625832,
    0.2310559 , 0.13364055, 0.43333333],
    [0.31428571, 0.43037975, 0.3362256 , 0.46107784, 0.5792277 ,
    0.50372671, 0.31797235, 0.33333333],
    [0.35714286, 0.39240506, 0.47722343, 0.52694611, 0.62849534,
    0.60714286, 0.35483871, 0.23333333],
    [0.11428571, 0.24050633, 0.21691974, 0.33532934, 0.38082557,
    0.34254658, 0.28110599, 0.4       ],
    [0.25714286, 0.36708861, 0.34924078, 0.80838323, 0.56724368,
    0.5173913 , 0.37788018, 0.23333333]])
```

```
[33]: import scipy
leng = feature_mtx.shape[0]
D = np.zeros([leng,leng])
for i in range(leng):
    for j in range(leng):
        D[i,j] = scipy.spatial.distance.euclidean(feature_mtx[i],
    ↪ feature_mtx[j])
```

```
[34]: Z = hierarchy.linkage(D, 'complete')
```

<ipython-input-34-f7fd5c287128>:1: ClusterWarning: scipy.cluster: The symmetric non-negative hollow observation matrix looks suspiciously like an uncondensed distance matrix


```
Z = hierarchy.linkage(D, 'complete')
```

```
[35]: #Hierarchical clustering does not require a pre-specified number of clusters.  
#However, in some applications we want a partition of disjoint clusters just as  
→in flat clustering.  
#So you can use a cutting line:
```

```
[36]: from scipy.cluster.hierarchy import fcluster  
max_d = 3  
clusters = fcluster(Z, max_d, criterion='distance')  
clusters
```

```
[36]: array([ 1,  5,  5,  6,  5,  4,  6,  5,  5,  5,  5,  5,  4,  4,  5,  1,  6,  
          5,  5,  5,  4,  2, 11,  6,  6,  5,  6,  5,  1,  6,  6, 10,  9,  8,  
          9,  3,  5,  1,  7,  6,  5,  3,  5,  3,  8,  7,  9,  2,  6,  6,  5,  
          4,  2,  1,  6,  5,  2,  7,  5,  5,  5,  4,  4,  3,  2,  6,  6,  5,  
          7,  4,  7,  6,  6,  5,  3,  5,  5,  6,  5,  4,  4,  1,  6,  5,  5,  
          5,  6,  4,  5,  4,  1,  6,  5,  6,  6,  5,  5,  5,  7,  7,  7,  2,  
          2,  1,  2,  6,  5,  1,  1,  1,  7,  8,  1,  1,  6,  1,  1],  
        dtype=int32)
```

```
[45]: import pylab  
fig = pylab.figure(figsize=(18,50))  
def llf(id):  
    return '%s %s %s' % (df['manufact'][id], df['model'][id],  
    →int(float(df['type'][id])))  
  
dendro = hierarchy.dendrogram(Z, leaf_label_func=llf, leaf_rotation=0,  
    →leaf_font_size =12, orientation = 'right')
```



```
[1]: #DBSCAN
```

```
[3]: import numpy as np
      from sklearn.cluster import DBSCAN
      from sklearn.preprocessing import StandardScaler
      from sklearn.datasets.samples_generator import make_blobs
      import matplotlib.pyplot as plt
      %matplotlib inline
```

C:\Users\Dell\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:143:
FutureWarning: The sklearn.datasets.samples_generator module is deprecated in
version 0.22 and will be removed in version 0.24. The corresponding classes /
functions should instead be imported from sklearn.datasets. Anything that cannot
be imported from sklearn.datasets is now part of the private API.

```
warnings.warn(message, FutureWarning)
```

```
[4]: def createDataPoints(centroidLocation, numSamples, clusterDeviation):
      # Create random data and store in feature matrix X and response vector y.
      X, y = make_blobs(n_samples=numSamples, centers=centroidLocation,
                        cluster_std=clusterDeviation)

      # Standardize features by removing the mean and scaling to unit variance
      X = StandardScaler().fit_transform(X)
      return X, y
```

```
[5]: X, y = createDataPoints([[4,3], [2,-1], [-1,4]] , 1500, 0.5)
```

```
[6]: db = DBSCAN(eps=0.3, min_samples=7)
```

```
[7]: db.fit(X)
```

```
[7]: DBSCAN(eps=0.3, min_samples=7)
```

```
[8]: labels=db.labels_
      labels
```

```
[8]: array([0, 1, 2, ..., 2, 0, 1], dtype=int64)
```

```
[10]: #distinguish outliers
      core_sample_mask = np.zeros_like(labels, dtype=bool)
      core_sample_mask[db.core_sample_indices_]=True
      core_sample_mask
```

```
[10]: array([ True,  True,  True, ...,  True,  True,  True])
```

```
[11]: # Number of clusters in labels, ignoring noise if present.  
n_clusters_ = len(set(labels)) - (1 if -1 in labels else 0)  
n_clusters_
```


```
[11]: 3
```

```
[12]: unique_label = set(labels)  
unique_label
```

```
[12]: {0, 1, 2}
```

```
[14]: # Create colors for the clusters.  
colors = plt.cm.Spectral(np.linspace(0, 1, len(unique_label)))
```

```
[18]: # Plot the points with colors  
for k, col in zip(unique_label, colors):  
    if k == -1:  
        # Black used for noise.  
        col = 'k'  
  
    class_member_mask = (labels == k)  
  
    # Plot the datapoints that are clustered  
    xy = X[class_member_mask & core_sample_mask]  
    plt.scatter(xy[:, 0], xy[:, 1], s=50, c=[col], marker='o', alpha=0.5)  
  
    # Plot the outliers  
    xy = X[class_member_mask & ~core_sample_mask]  
    plt.scatter(xy[:, 0], xy[:, 1], s=50, c=[col], marker='o', alpha=0.5)
```

```
[19]: #using dataset
```

```
[20]: import pandas as pd
import numpy as np
```

```
[23]: pdf = pd.read_csv('weather-stations20140101-20141231.csv')
```

```
[24]: pdf.head()
```

```
[24]:
```

	Stn_Name	Lat	Long	Prov	Tm	DwTm	D	Tx	DwTx	\
0	CHEMAINUS	48.935	-123.742	BC	8.2	0.0	NaN	13.5	0.0	
1	COWICHAN LAKE FORESTRY	48.824	-124.133	BC	7.0	0.0	3.0	15.0	0.0	
2	LAKE COWICHAN	48.829	-124.052	BC	6.8	13.0	2.8	16.0	9.0	
3	DISCOVERY ISLAND	48.425	-123.226	BC	NaN	NaN	NaN	12.5	0.0	
4	DUNCAN KELVIN CREEK	48.735	-123.728	BC	7.7	2.0	3.4	14.5	2.0	

	Tn	...	DwP	P%N	S_G	Pd	BS	DwBS	BS%	HDD	CDD	Stn_No
0	1.0	...	0.0	NaN	0.0	12.0	NaN	NaN	NaN	273.3	0.0	1011500
1	-3.0	...	0.0	104.0	0.0	12.0	NaN	NaN	NaN	307.0	0.0	1012040
2	-2.5	...	9.0	NaN	NaN	11.0	NaN	NaN	NaN	168.1	0.0	1012055
3	NaN	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	1012475
4	-1.0	...	2.0	NaN	NaN	11.0	NaN	NaN	NaN	267.7	0.0	1012573

```
[5 rows x 25 columns]
```

```
[25]: pdf = pdf[pd.notnull(pdf["Tm"])]
pdf = pdf.reset_index(drop=True)
pdf.head(5)
```

```
[25]:
```

	Stn_Name	Lat	Long	Prov	Tm	DwTm	D	Tx	DwTx	\
0	CHEMAINUS	48.935	-123.742	BC	8.2	0.0	NaN	13.5	0.0	
1	COWICHAN LAKE FORESTRY	48.824	-124.133	BC	7.0	0.0	3.0	15.0	0.0	
2	LAKE COWICHAN	48.829	-124.052	BC	6.8	13.0	2.8	16.0	9.0	
3	DUNCAN KELVIN CREEK	48.735	-123.728	BC	7.7	2.0	3.4	14.5	2.0	
4	ESQUIMALT HARBOUR	48.432	-123.439	BC	8.8	0.0	NaN	13.1	0.0	

	Tn	...	DwP	P%N	S_G	Pd	BS	DwBS	BS%	HDD	CDD	Stn_No
0	1.0	...	0.0	NaN	0.0	12.0	NaN	NaN	NaN	273.3	0.0	1011500
1	-3.0	...	0.0	104.0	0.0	12.0	NaN	NaN	NaN	307.0	0.0	1012040
2	-2.5	...	9.0	NaN	NaN	11.0	NaN	NaN	NaN	168.1	0.0	1012055
3	-1.0	...	2.0	NaN	NaN	11.0	NaN	NaN	NaN	267.7	0.0	1012573
4	1.9	...	8.0	NaN	NaN	12.0	NaN	NaN	NaN	258.6	0.0	1012710

[5 rows x 25 columns]

```
[26]: from mpl_toolkits.basemap import Basemap
import matplotlib.pyplot as plt
from pylab import rcParams
%matplotlib inline
rcParams['figure.figsize'] = (14,10)

llon=-140
ulon=-50
llat=40
ulat=65

pdf = pdf[(pdf['Long'] > llon) & (pdf['Long'] < ulon) & (pdf['Lat'] > llat)
->&(pdf['Lat'] < ulat)]

my_map = Basemap(projection='merc',
                  resolution = 'l', area_thresh = 1000.0,
                  llcrnrlon=llon, llcrnrlat=llat, #min longitude (llcrnrlon) and
->latitude (llcrnrlat)
                  urcrnrlon=ulon, urcrnrlat=ulat) #max longitude (urcrnrlon) and
->latitude (urcrnrlat)

my_map.drawcoastlines()
my_map.drawcountries()
# my_map.drawmapboundary()
my_map.fillcontinents(color = 'white', alpha = 0.3)
my_map.shadedrelief()
```

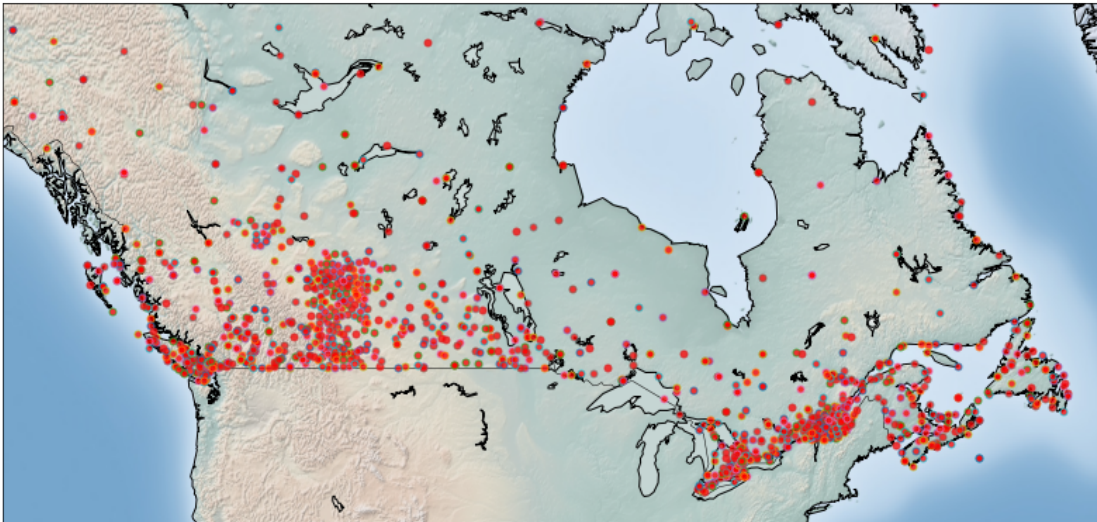
```

# To collect data based on stations

xs,ys = my_map(np.asarray(pdf.Long), np.asarray(pdf.Lat))
pdf['xm']= xs.tolist()
pdf['ym']=ys.tolist()

#Visualization1
for index,row in pdf.iterrows():
    # x,y = my_map(row.Long, row.Lat)
    my_map.plot(row.xm, row.ym,markerfacecolor =([1,0,0]), marker='o',
    ↪markersize= 5, alpha = 0.75)
#plt.text(x,y,stn)
plt.show()

```



```

[27]: from sklearn.cluster import DBSCAN
import sklearn.utils
from sklearn.preprocessing import StandardScaler
sklearn.utils.check_random_state(1000)
Clus_dataSet = pdf[['xm','ym']]
Clus_dataSet = np.nan_to_num(Clus_dataSet)
Clus_dataSet = StandardScaler().fit_transform(Clus_dataSet)

# Compute DBSCAN
db = DBSCAN(eps=0.15, min_samples=10).fit(Clus_dataSet)
core_samples_mask = np.zeros_like(db.labels_, dtype=bool)
core_samples_mask[db.core_sample_indices_] = True
labels = db.labels_
pdf["Clus_Db"]=labels

```

```
realClusterNum=len(set(labels)) - (1 if -1 in labels else 0)
clusterNum = len(set(labels))
```

```
# A sample of clusters
pdf[["Stn_Name", "Tx", "Tm", "Clus_Db"]].head(5)
```

```
[27]:
```

	Stn_Name	Tx	Tm	Clus_Db
0	CHEMAINUS	13.5	8.2	0
1	COWICHAN LAKE FORESTRY	15.0	7.0	0
2	LAKE COWICHAN	16.0	6.8	0
3	DUNCAN KELVIN CREEK	14.5	7.7	0
4	ESQUIMALT HARBOUR	13.1	8.8	0

```
[28]: set(labels)
```

```
[28]: {-1, 0, 1, 2, 3, 4}
```

```
[29]: #visualization of clusters based on location
```

```
[30]: from mpl_toolkits.basemap import Basemap
import matplotlib.pyplot as plt
from pylab import rcParams
%matplotlib inline
rcParams['figure.figsize'] = (14,10)

my_map = Basemap(projection='merc',
                  resolution = 'l', area_thresh = 1000.0,
                  llcrnrlon=llon, llcrnrlat=llat, #min longitude (llcrnrlon) and
↳ latitude (llcrnrlat)
                  urcrnrlon=ulon, urcrnrlat=ulat) #max longitude (urcrnrlon) and
↳ latitude (urcrnrlat)

my_map.drawcoastlines()
my_map.drawcountries()
#my_map.drawmapboundary()
my_map.fillcontinents(color = 'white', alpha = 0.3)
my_map.shadedrelief()

# To create a color map
colors = plt.get_cmap('jet')(np.linspace(0.0, 1.0, clusterNum))

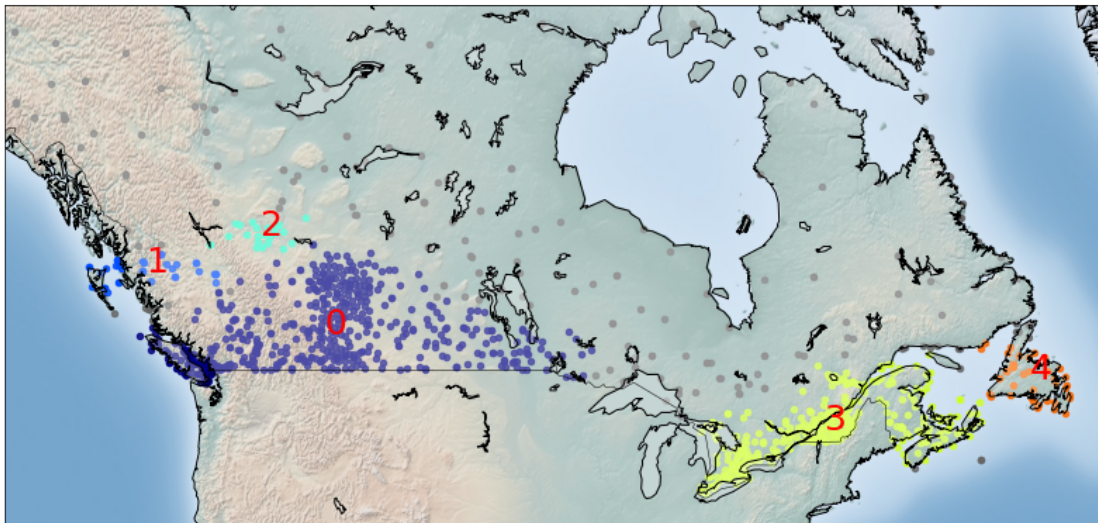
#Visualization1
for clust_number in set(labels):
```

```

c=([0.4,0.4,0.4]) if clust_number == -1 else colors[np.int(clust_number)])
clust_set = pdf[pdf.Clus_Db == clust_number]
my_map.scatter(clust_set.xm, clust_set.ym, color =c, marker='o', s= 20,
↳alpha = 0.85)
if clust_number != -1:
    cenx=np.mean(clust_set.xm)
    ceny=np.mean(clust_set.ym)
    plt.text(cenx,ceny,str(clust_number), fontsize=25, color='red',)
    print ("Cluster "+str(clust_number)+', Avg Temp: '+ str(np.
↳mean(clust_set.Tm)))

```

Cluster 0, Avg Temp: -5.538747553816051
Cluster 1, Avg Temp: 1.9526315789473685
Cluster 2, Avg Temp: -9.195652173913045
Cluster 3, Avg Temp: -15.300833333333333
Cluster 4, Avg Temp: -7.769047619047619



[]: *#Clustering of stations based on their location, mean, max, and min Temperature*

```

[31]: from sklearn.cluster import DBSCAN
import sklearn.utils
from sklearn.preprocessing import StandardScaler
sklearn.utils.check_random_state(1000)
Clus_dataSet = pdf[['xm', 'ym', 'Tx', 'Tm', 'Tn']]
Clus_dataSet = np.nan_to_num(Clus_dataSet)
Clus_dataSet = StandardScaler().fit_transform(Clus_dataSet)

# Compute DBSCAN
db = DBSCAN(eps=0.3, min_samples=10).fit(Clus_dataSet)

```

```

core_samples_mask = np.zeros_like(db.labels_, dtype=bool)
core_samples_mask[db.core_sample_indices_] = True
labels = db.labels_
pdf["Clus_Db"]=labels

realClusterNum=len(set(labels)) - (1 if -1 in labels else 0)
clusterNum = len(set(labels))

# A sample of clusters
pdf[["Stn_Name", "Tx", "Tm", "Clus_Db"]].head(5)

```

```

[31]:
      Stn_Name    Tx    Tm  Clus_Db
0      CHEMAINUS  13.5  8.2         0
1  COWICHAN LAKE FORESTRY  15.0  7.0         0
2      LAKE COWICHAN  16.0  6.8         0
3  DUNCAN KELVIN CREEK  14.5  7.7         0
4  ESQUIMALT HARBOUR   13.1  8.8         0

```

```

[32]: #visualization based on location and temperature

```

```

[33]: from mpl_toolkits.basemap import Basemap
import matplotlib.pyplot as plt
from pylab import rcParams
%matplotlib inline
rcParams['figure.figsize'] = (14,10)

my_map = Basemap(projection='merc',
                  resolution = 'l', area_thresh = 1000.0,
                  llcrnrlon=llon, llcrnrlat=llat, #min longitude (llcrnrlon) and
↳latitude (llcrnrlat)
                  urcrnrlon=ulon, urcrnrlat=ulat) #max longitude (urcrnrlon) and
↳latitude (urcrnrlat)

my_map.drawcoastlines()
my_map.drawcountries()
#my_map.drawmapboundary()
my_map.fillcontinents(color = 'white', alpha = 0.3)
my_map.shadedrelief()

# To create a color map
colors = plt.get_cmap('jet')(np.linspace(0.0, 1.0, clusterNum))

#Visualization1
for clust_number in set(labels):

```



```

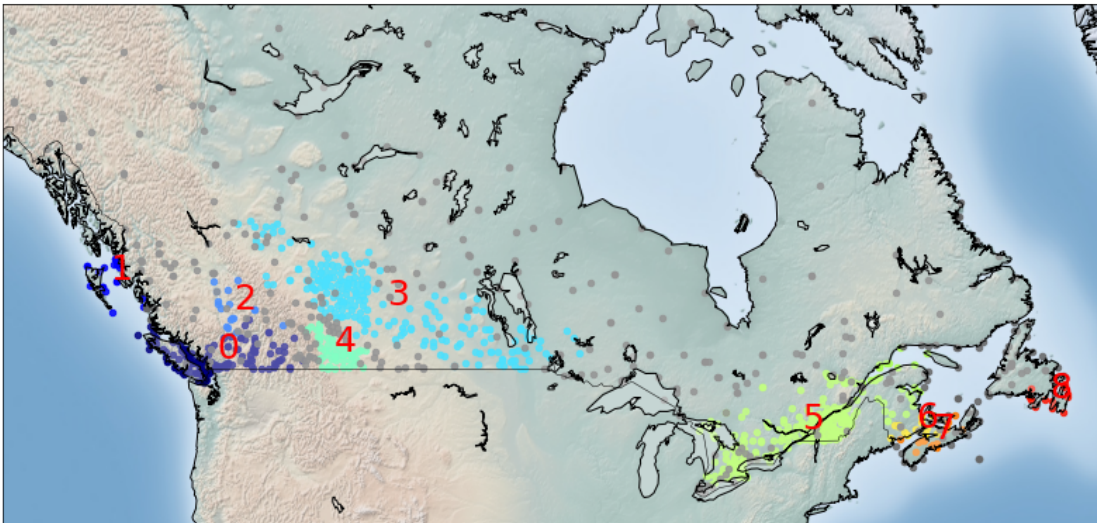
c=([0.4,0.4,0.4]) if clust_number == -1 else colors[np.int(clust_number)])
clust_set = pdf[pdf.Clus_Db == clust_number]
my_map.scatter(clust_set.xm, clust_set.ym, color =c, marker='o', s= 20,
↳alpha = 0.85)
if clust_number != -1:
    cenx=np.mean(clust_set.xm)
    ceny=np.mean(clust_set.ym)
    plt.text(cenx,ceny,str(clust_number), fontsize=25, color='red',)
    print ("Cluster "+str(clust_number)+', Avg Temp: '+ str(np.
↳mean(clust_set.Tm)))

```

```

Cluster 0, Avg Temp: 6.2211920529801334
Cluster 1, Avg Temp: 6.7900000000000001
Cluster 2, Avg Temp: -0.49411764705882355
Cluster 3, Avg Temp: -13.877209302325586
Cluster 4, Avg Temp: -4.186274509803922
Cluster 5, Avg Temp: -16.301503759398482
Cluster 6, Avg Temp: -13.599999999999998
Cluster 7, Avg Temp: -9.753333333333334
Cluster 8, Avg Temp: -4.258333333333334

```



[]: